

EP54: Cache Systems Every Developer Should Know



ALEX XU

APR 8, 2023

143

4

4

Share

...

This week's system design refresher:

- Cache Systems Every Developer Should Know (Youtube video)
- Serverless DB
- TCP vs. UDP
- Batch v.s. Stream Processing
- Job openings

Become a better tech leader in 10 mins a week (Sponsored)

Become a better tech leader with Refactoring

Every week, 30K+ engineers and managers at big tech and high-growth startups read Refactoring to get better at their job

	DEV EFFORT	CHANGE CONTRACT	SWAP IMPLEMENTATION	ROI
UNIT TEST	LOW	OPTIONAL	RARELY	MEDIUM
INTEGRATION TEST	MEDIUM	RARELY	OPTIONAL	HIGH
END TO END TEST	HIGH	OPTIONAL	OPTIONAL	MEDIUM

ISSUE SEVERITY
PEOPLE LEAVE
TIME (WEEKS)
BI-WEEKLY 1:1s
WEEKLY 1:1s

CROSS-TEAM
STAFF ENGINEER
ENGINEER
EM / PM
SINGLE TEAM
TPM
PLANNING + PROJ. MGMT

Refactoring is a weekly newsletter about writing great software and working well together. Its articles are read every week by 30,000+ subscribers, including engineers

and managers at companies like Google, Meta, and Amazon.

Subscribe to:

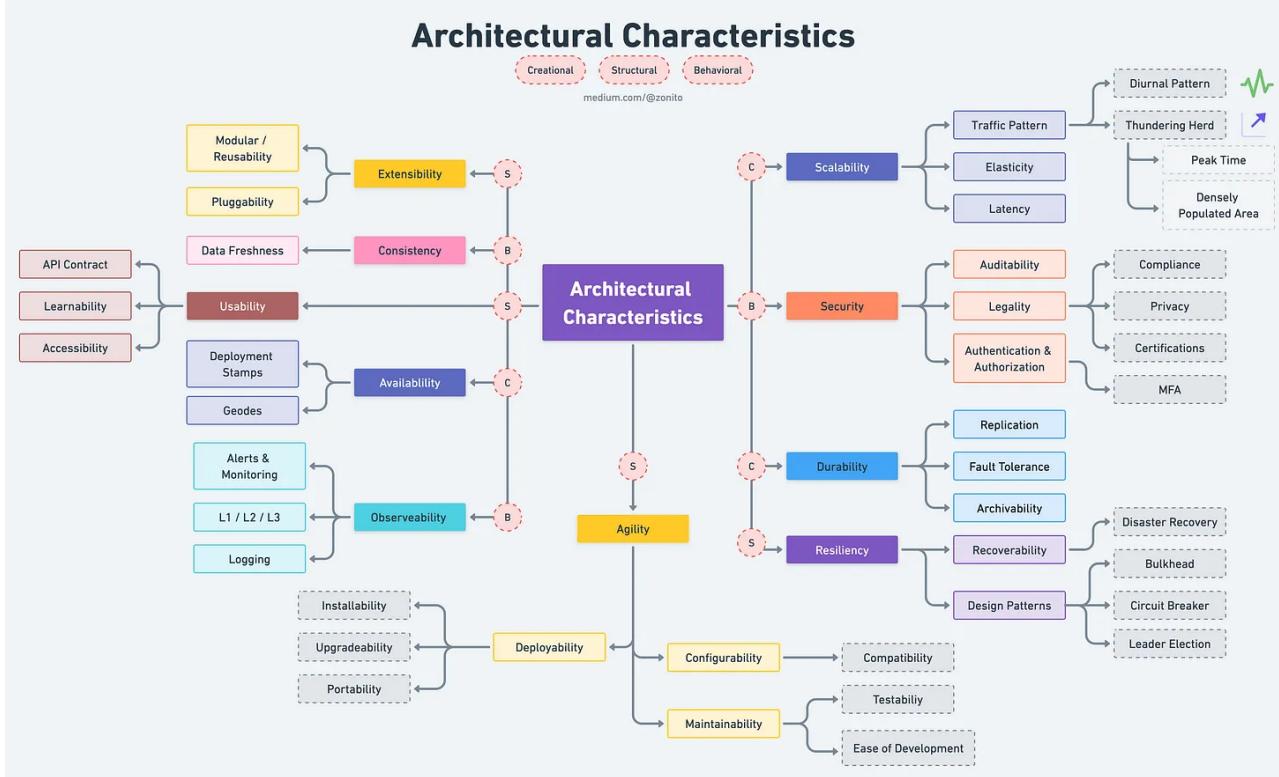
- Receive a **new essay** every Thursday with practical advice about your work.
- Access a **curated library** of 130+ original essays and 250+ resources.
- Join a **private community** of tech leaders, founders, and engineers.
- Get **\$50K+** discounts on popular dev tools.

Cache Systems Every Developer Should Know

Cache Systems Every Developer Should Know



Top 10 Architecture Characteristics / Non-Functional Requirements with Cheatsheet



Did we miss anything? If yes, Please help to enrich this article by sharing your thoughts in the comments.

Written by [Love Sharma](#), our guest author. We're constantly seeking valuable content, so if you'd like to contribute to our platform or have any previously published content you'd like us to share, please feel free to drop us a message.

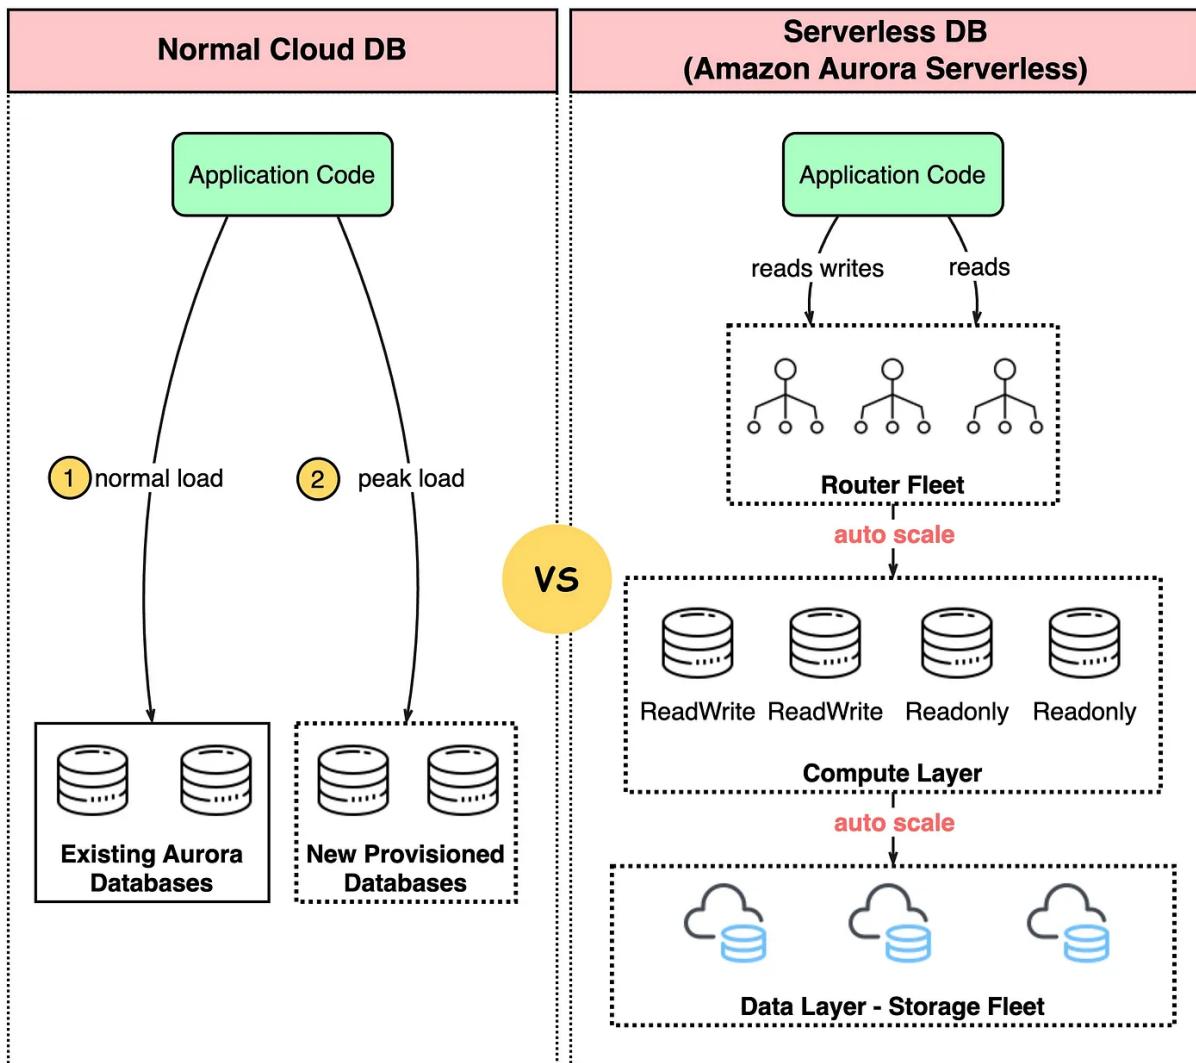
What is Serverless DB

Are serverless databases the future? How do serverless databases differ from traditional cloud databases?

Amazon Aurora Serverless, depicted in the diagram below, is a configuration that is auto-scaling and available on-demand for Amazon Aurora.

What is Serverless DB?

 blog.bytebytego.com



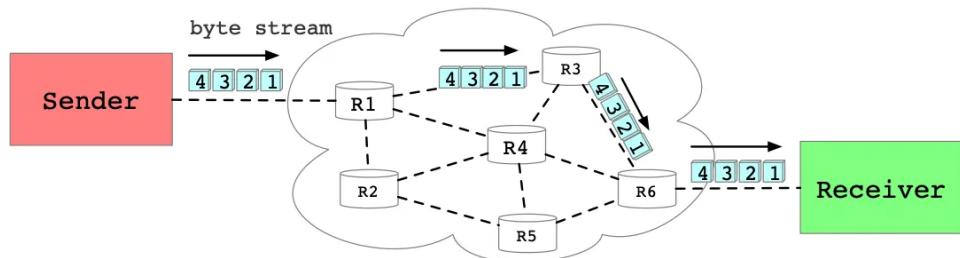
- Aurora Serverless has the ability to scale capacity automatically up or down as per business requirements. For example, an eCommerce website preparing for a major promotion can scale the load to multiple databases within a few milliseconds. In comparison to regular cloud databases, which necessitate the provision and administration of database instances, Aurora Serverless can automatically start up and shut down.
- By decoupling the compute layer from the data storage layer, Aurora Serverless is able to charge fees in a more precise manner. Additionally, Aurora Serverless can be a combination of provisioned and serverless instances, enabling existing provisioned databases to become a part of the serverless pool.

Over to you: Have you used a serverless DB? Does it save cost?

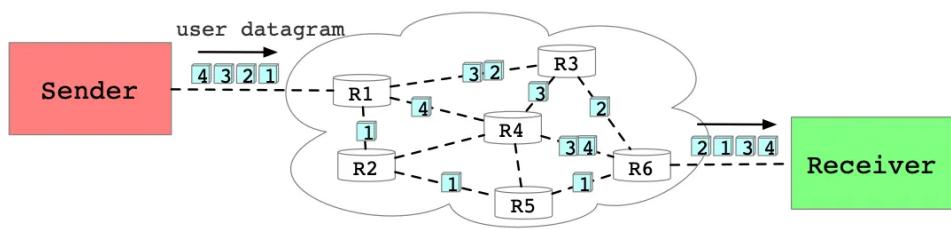
TCP vs. UDP: 7 Differences You Should Know

1. Connection-oriented vs. connectionless

TCP vs. UDP: 7 Differences You Should Know



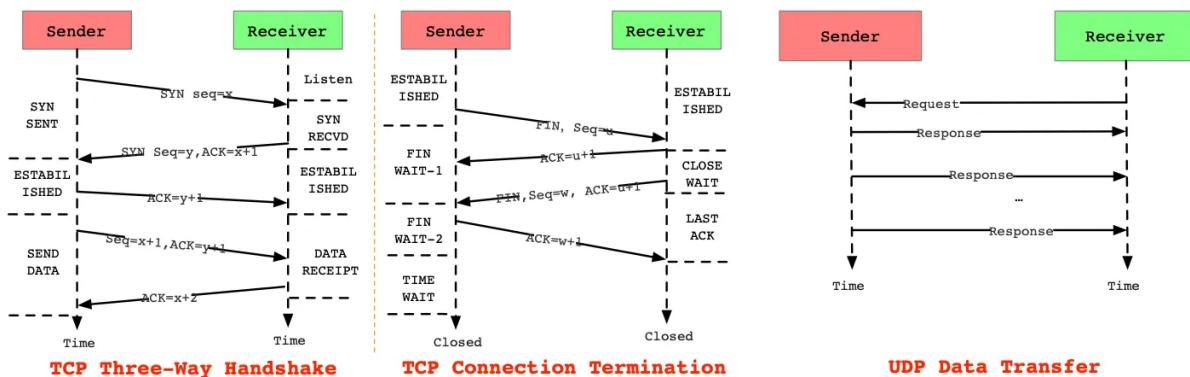
TCP connection-oriented
Data arrives in-order



UDP connectionless
Data could be out of order

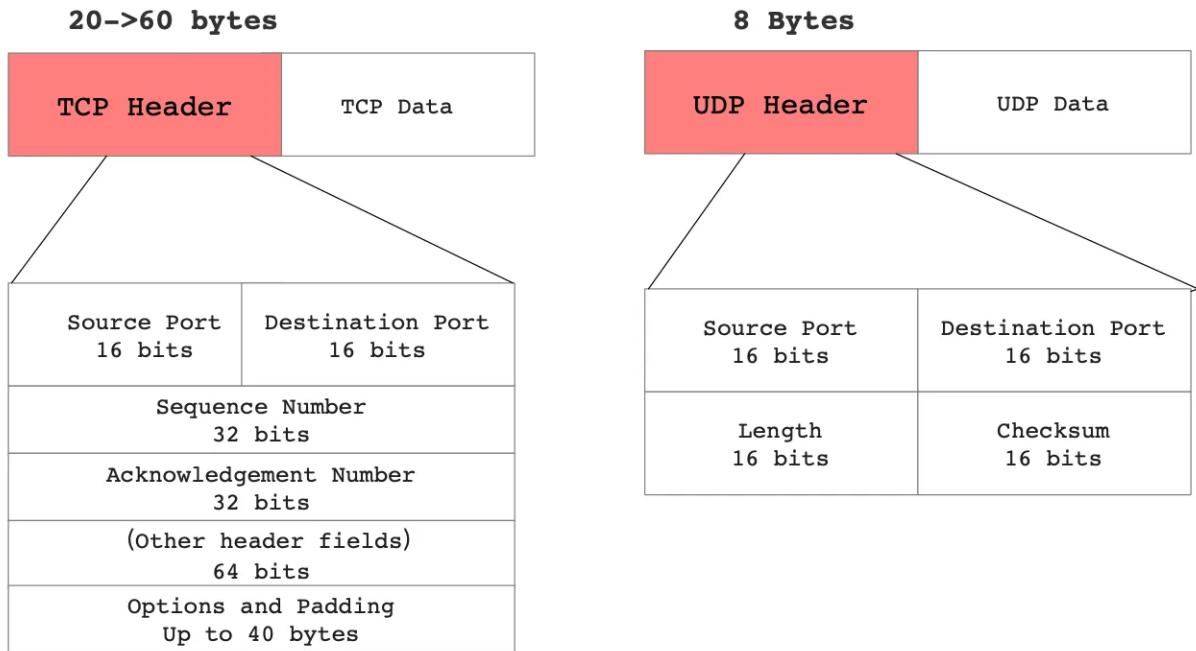
2. Three-way handshake vs. No handshake

Three-way handshake vs. No Handshake



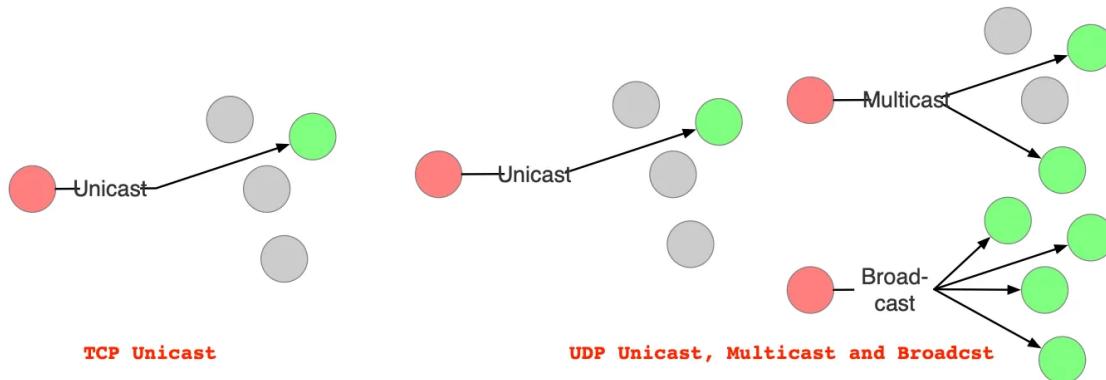
3. Header (20 bytes) vs. (8 bytes)

header (20 bytes) vs. header (8 bytes)



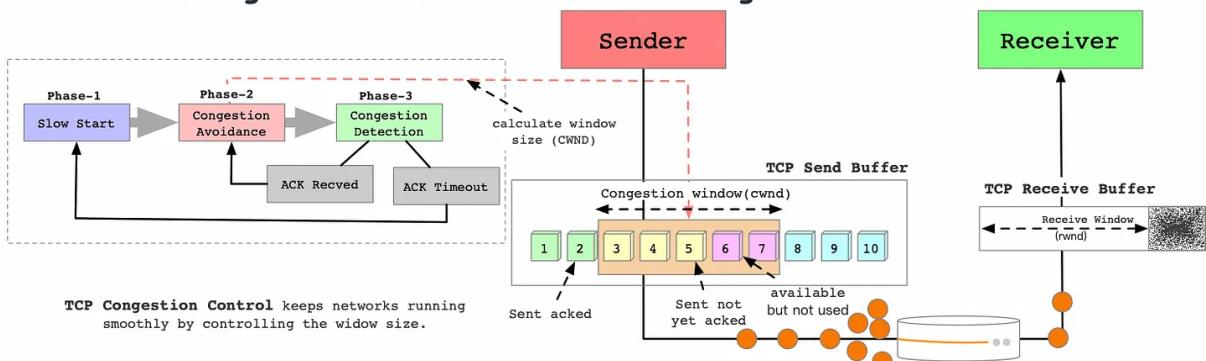
4. Point-to-point vs. Unicast & Multicast & Broadcast

Point to point vs. Unicast & Multicast & Broadcast



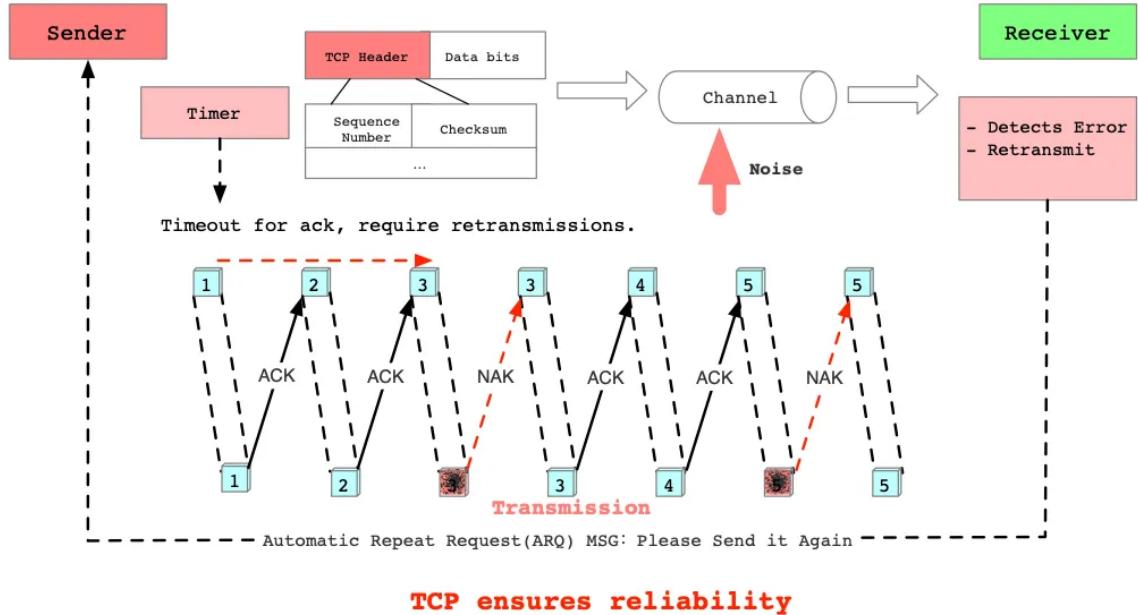
5. Congestion control vs. no congestion control

Congestion control vs. No congestion control



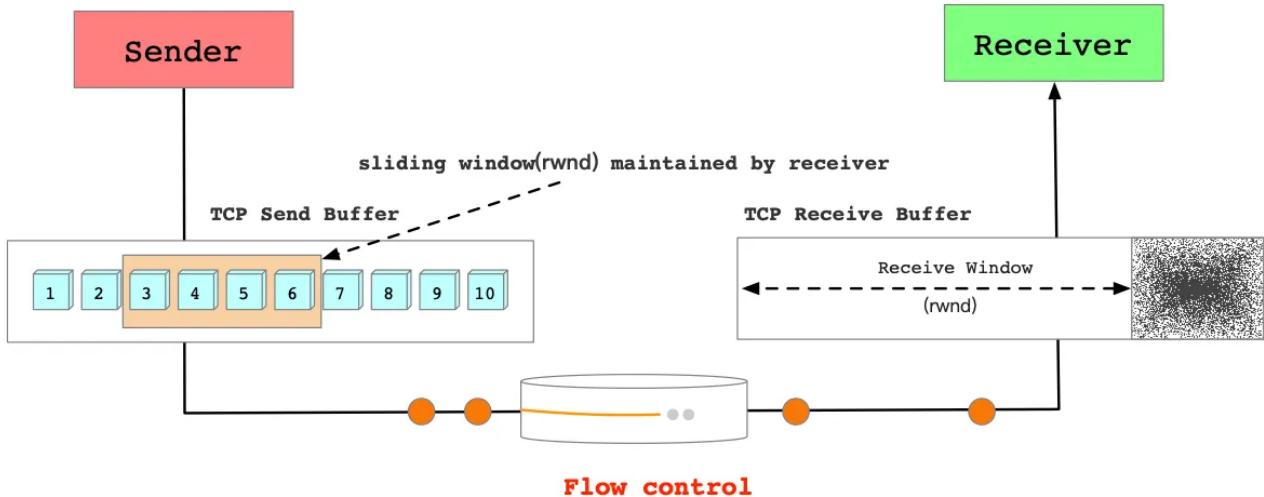
6. Reliable vs. lossy

Reliable vs. Lossy



7. Flow control vs. no flow control

Flow control vs. No flow Control

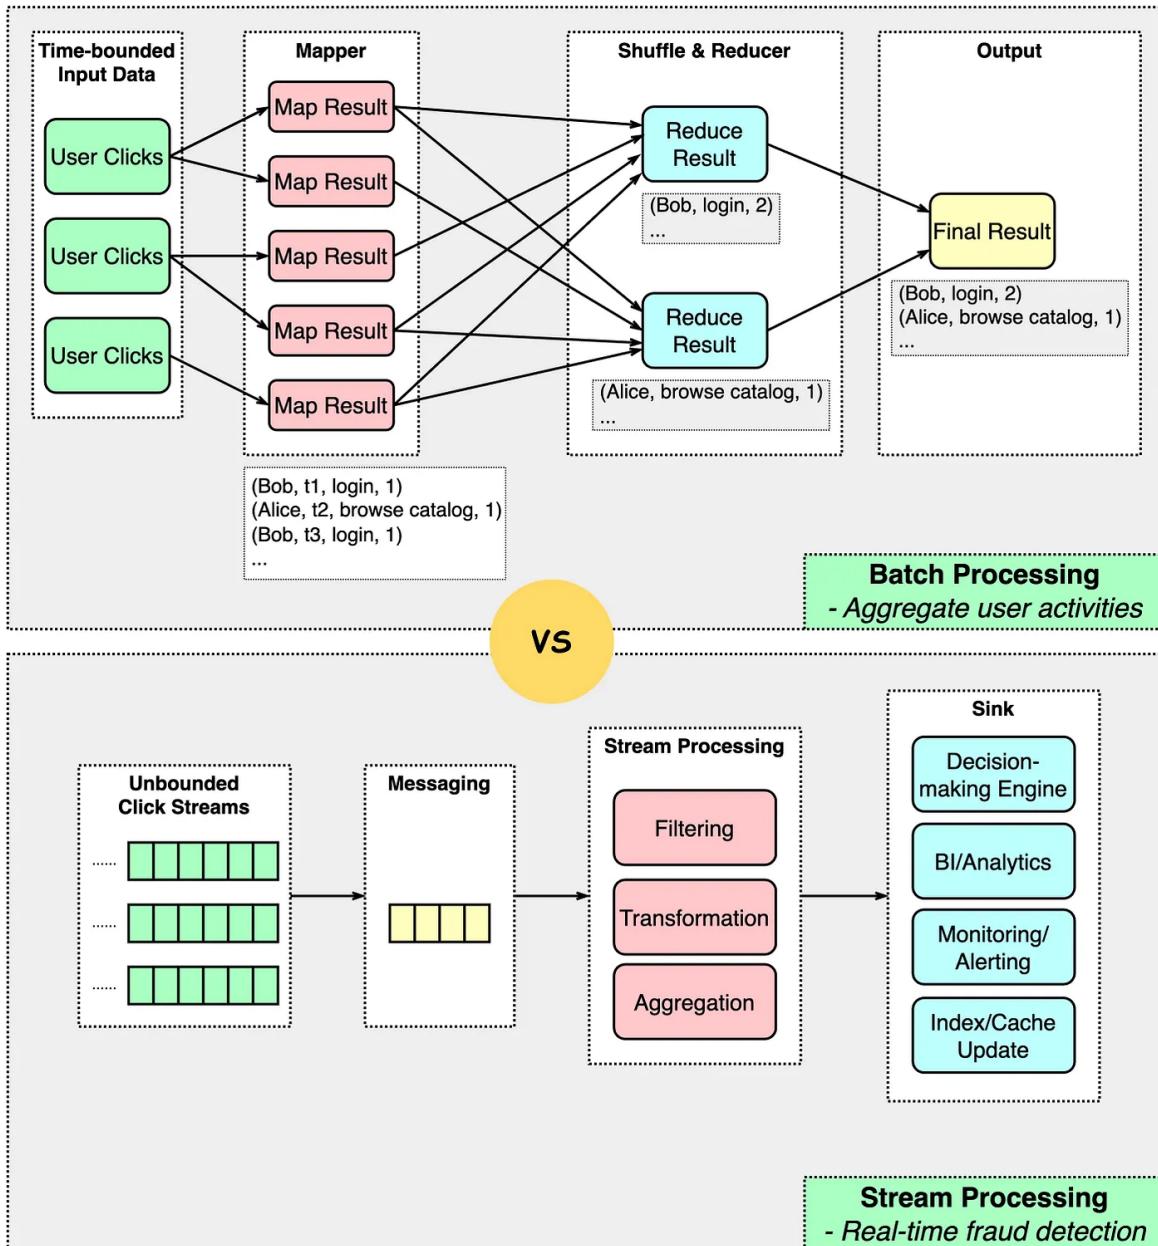


Batch v.s. Stream Processing

- Batch Processing: We aggregate user click activities at end of the day.
- Stream Processing: We detect potential frauds with the user click streams in real-time.

Batch v.s. Stream Processing

 blog.bytebytego.com



Both processing models are used in big data processing. The major differences are:

1. Input

Batch processing works on time-bounded data, which means there is an end to the input data.

Stream processing works on data streams, which doesn't have a boundary.

2. Timeliness

Batch processing is used in scenarios where the data doesn't need to be

processed in real-time.

Stream processing can produce processing results as the data is generated.

3. Output

Batch processing usually generates one-off results, for example, reports.

Stream processing's outputs can pipe into fraud decision-making engines, monitoring tools, analytics tools, or index/cache updaters.

4. Fault tolerance

Batch processing tolerates faults better as the batch can be replayed on a fixed set of input data.

Stream processing is more challenging as the input data keeps flowing in. There are some approaches to solve this:

- a) Microbatching which splits the data stream into smaller blocks (used in Spark);
- b) Checkpoint which generates a mark every few seconds to roll back to (used in Flink).

Over to you: Have you worked on stream processing systems?

Join the ByteByteGo Talent Collective

If you're looking for a new gig, [join the collective](#) for customized job offerings from selected companies. Public or anonymous options are available. Leave anytime.



ByteByteGo

Curated By  Alex Xu

Bringing together top-performing talents and leading tech companies.

Engineering Manager
Brilliant
 Senior-Expert
San Francisco, Remote (North America)

Head of Infrastructure and Risk
X1
 Expert
Remote (US)

Software Engineer, Back End
X1
 All Levels
Remote (US)

Fullstack Engineer
Equi
 Senior-Expert
New York

Alex Xu is on [Pallet](#)

If you're **hiring**, [join the ByteByteGo Talent Collective](#) to start getting bi-monthly drops of world-class hand-curated engineers who are open to new opportunities.

Featured job openings

X1 Card: [Engineering Leader - Card Platform](#) (United States, remote)



143 Likes · 4 Restacks

4 Comments



Write a comment...



Jakub Jablonski Writes Jakub's Substack Apr 10

Regarding architecture characteristics, I have started collecting a library of those for myself, see <https://jakubjablonski2-tomtom.github.io/architecture-characteristics/>.

LIKE (2) REPLY SHARE

...



John Stanhope Writes John's Substack Apr 15

I would also consider performance (resource efficiency) and latency as an architectural characteristics.

LIKE REPLY SHARE

...

2 more comments...