

# Metric monitoring



ALEX XU

JAN 12, 2022



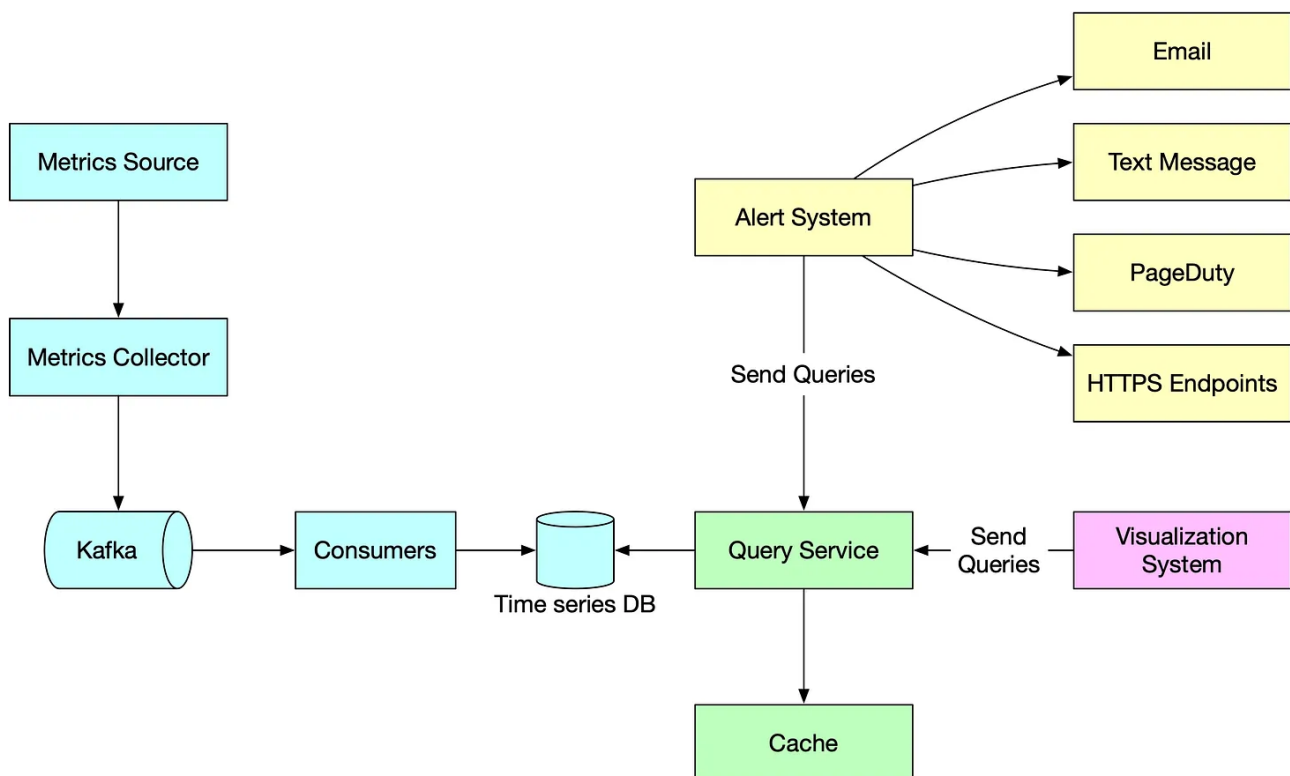
28



Share



A well-designed **metric monitoring** and alerting system plays a key role in providing clear visibility into the health of the infrastructure to ensure high availability and reliability. The diagram below explains how it works at a high level.



**Metrics source:** This can be application servers, SQL databases, message queues, etc.

**Metrics collector:** It gathers metrics data and writes data into the time-series database.

**Time-series database:** This stores metrics data as time series. It usually provides a custom query interface for analyzing and summarizing a large amount of time-series data. It maintains indexes on labels to facilitate the fast lookup of time-series data by labels.

**Kafka:** Kafka is used as a highly reliable and scalable distributed messaging platform. It decouples the data collection and data processing services from each other.

**Consumers:** Consumers or streaming processing services such as Apache Storm, Flink and Spark, process and push data to the time-series database.

**Query service:** The query service makes it easy to query and retrieve data from the time-series database. This should be a very thin wrapper if we choose a good time-series database. It could also be entirely replaced by the time-series database's own query interface.

**Alerting system:** This sends alert notifications to various alerting destinations.

**Visualization system:** This shows metrics in the form of various graphs/charts.

If you enjoyed this post, you might like our system design interview books as well.

SDI-vol1: <https://amzn.to/3tK0qQn>

SDI-vol2: <https://amzn.to/37ZisW9>




28 Likes

## Comments



Write a comment...



© 2023 ByteByteGo · [Privacy](#) · [Terms](#) · [Collection notice](#)  
[Substack](#) is the home for great writing